

Day 3: Portfolio Credit Risk, Marginal Survival & Dependence Modeling Report

Course: Machine Learning in Industry · Cardo AI / MSCA Digital Doctoral Network

Module: Copula-Based Portfolio Risk · Marginal & Dependence Modeling · Model Validation

AI Disclosure Statement: Large Language Model (LLM) tooling was utilized for code scaffolding, vectorization debugging, and report structuring; all empirical estimations, latent parameter bounds, and simulation results were verified and analyzed independently.

1. Executive Summary & Plain-English Intuition

When evaluating an Asset-Backed Security (ABS) collateralized by personal consumer loans, the single most critical question is not just how often loans default on average, but whether they default at the same time during an economic downturn.

In reality, a macroeconomic recession hits everyone simultaneously. When the economy crashes, borrowers lose jobs together and default together.

This project conducted an independent audit of a raw loan portfolio:

- Cleaned and reconstructed 287,819 legitimate loans from a heavily corrupted production tape.
- Estimated the individual default timing over a loan's lifespan using parametric survival analysis.
- Quantified the hidden systematic co-movement (asset correlation) linking obligors via a latent-factor Gaussian/Vasicek copula.
- Simulated portfolio tail losses, proving that treating loans as independent severely understates tail risk by 28.7x, requiring a 32% credit enhancement to protect senior investors.

A critical post-estimation review reveals that the fitted asset correlation ($\rho = 0.50$) hit the upper search boundary of the optimization routine, indicating an estimation artifact rather than pure macroeconomic co-movement. The aggregation by origination year (vintage) conflates portfolio seasoning and observation-window right-censoring with systemic calendar-year shocks. In regulatory consumer credit frameworks (such as Basel IRB), retail asset correlations typically fluctuate between 3% and 12%.

Consequently, while the 28.7x capital understatement under independence demonstrates the severe vulnerability of assuming uncorrelated risks, the 1.41 million Euro tail loss should be interpreted as an extreme macro-stress scenario rather than a

through-the-cycle baseline. Future iterations must condition default timing on calendar-year macroeconomic covariates to decouple structural seasoning from systemic crisis events.

2. Part 1: Data Audit, Pipeline Cleaning & Exposure at Default (EAD)

The raw dataset ([loan_tape.csv](#)) arrived as an uncurated operational dump of 308,222 records characterized by significant data corruption:

- Origination Date Recovery: 10,317 records lacked an origination date. Using contractual maturity dates and term lengths ($t = \text{maturity} - \text{term}$), these missing dates were mathematically recovered without discarding loans.
- Corrupted Event Dates: Found 6,053 instances where a default date occurred *before* the loan was originated, and 4,468 prepayment dates preceding origination. These impossible records were sanitized to [NaN](#).
- Typo Standardization & Outlier Removal: Standardized messy categorical text ([Northe@st](#), [Frnch](#), [XYZ](#), [ITALIAN](#)) into clean amortization types ([bullet](#), [italian](#), [french](#)). Filtered out negative interest rates, rates exceeding 40%, and nonsensical loan sizes (including records exceeding €50M).
- Clean Population: Retained 287,819 valid obligors (96.23% retention rate).
 - *Censored (Still Active/Matured without Event)*: 79.54% (228,929 loans)
 - *Prepaid Early*: 10.88% (31,312 loans)
 - *Defaulted*: 9.58% (27,578 loans)
- Amortization Engine & EAD: Reconstructed closed-form monthly amortization schedules for French (constant annuity), Italian (constant linear principal amortization), and Bullet (interest-only with lump-sum balance at maturity) loans. Calculated exact outstanding Exposure at Default (EAD) as of the observation cut-off date ([obs_end_date](#)).

3. Part 2: Marginal Survival Modeling & Competing Risks

To project when defaults happen, the time scale was defined as loan age in months since origination (t), rather than calendar time, ensuring loans issued in different years are compared fairly along their repayment lifecycle.

- Competing Risks Recognition: Default and prepayment are mutually exclusive absorbing states. A borrower who refinances or pays off a loan early can no longer default. Treating prepayment as neutral right-censoring would distort default risk upward.
- Distribution Selection via Maximum Likelihood Estimation (MLE):
 - Log-Normal Distribution: $\mu = 1.937$, $\sigma^2 = 431.49$ | AIC: 65,576.5 (Selected Champion)
 - Weibull Distribution: $c = 0.970$, $\lambda = 395.27$ | AIC: 66,030.5

- Exponential Distribution: $\text{scale} = 371.48$ | AIC: 66,036.3
- *Economic Rationale:* Consumer credit risk exhibits a non-monotonic hazard curve: default risk is low during initial months, peaks midway as borrowers encounter liquidity shocks, and tapers off as the balance amortizes. The Log-Normal model captures this seasoning curve with the lowest AIC.

4. Part 3: Latent Factor Architecture & Dependence Modeling (Copula)

Borrower credit quality is driven by both an idiosyncratic personal shock and a shared macroeconomic factor. We modeled this relationship using the standard single-factor Vasicek/Gaussian copula framework:

$$X_i = \sqrt{\rho} Z + \sqrt{1 - \rho} \epsilon_i$$

Where $Z \sim \mathcal{N}(0, 1)$ is the macroeconomic market shock, $\epsilon_i \sim \mathcal{N}(0, 1)$ is borrower-specific risk, and ρ represents the asset correlation.

- Fitted Asset Correlation: Estimating cross-cohort default volatility using Maximum Likelihood yielded an asset correlation of $\rho = 0.50$. This indicates that macro shocks drive 50% of asset variance across the loan book.
- Tail Dependence Analysis:
 - The standard Gaussian copula assumes zero asymptotic tail dependence ($\lambda_L = 0$), implying extreme joint defaults almost never cluster together.
 - A Student- t structure with 4 degrees of freedom ($\nu = 4$) yields an implied lower tail dependence of $\lambda_L = 0.2532$. This demonstrates that real-world economic crises generate severe simultaneous default clustering far beyond Gaussian assumptions.

5. Part 4: The Regional Diversification Myth

The loan originator claimed the pool is "safely diversified across multiple geographic regions". Our empirical audit directly refutes this claim:

1. Regional Baseline Divergence: Baseline default rates are structurally uneven (e.g., Northeast at 16.6% vs Southeast at 6.2%).
2. Failure of Diversification in Recessions: Because obligors share a heavy asset correlation ($\rho = 0.50$) to the global economy, local geographic diversification breaks down during severe downturns. All regional pools load onto the same systemic shock simultaneously.

6. Part 5: Monte Carlo Portfolio Loss Simulation & Risk Metrics

Using the fitted Log-Normal marginals, 10,000 Monte Carlo paths were simulated over a 12-month forward horizon with an industry-standard unsecured Loss Given Default (LGD) = 45\%).

Portfolio Tail Loss Comparison Table

Metric	Naive Independence ($\rho=0.00$)	Fitted Copula ($\rho=0.50$)	Difference / Impact
Expected Loss (EL)	€66,012.54	€66,669.68	Comparable baseline
Value at Risk (VaR 95%)	€90,216.34	€346,961.50	3.8x increase
Value at Risk (VaR 99%)	€101,427.91	€790,357.68	7.8x increase
Value at Risk (VaR 99.9%)	€112,906.96	€1,411,391.68	12.5x increase
Expected Shortfall (ES 99.9%)	€118,866.40	€1,611,788.93	13.6x increase
Economic Capital ($\text{VaR}_{99.9} - \text{EL}$)	€46,894.43	€1,344,722.00	28.7x capital shortfall

Portfolio Loss Rate at 99.9%	2.48%	31.02%	Severe tail exposure
------------------------------	-------	--------	----------------------

- Key Takeaway: While the Expected Loss (average year) is virtually identical (~€66k), the 1-in-1,000-year crisis scenario (VaR 99.9%) differs radically. Assuming independence suggests a maximum loss of €113k, whereas correlated defaults destroy €1.41M.

7. Part 6: Model Limitations & Desk Recommendations

- Model Risks: Static copulas overlook dynamic feedback effects, deteriorating recovery rates (downturn LGD), and prepayment slowdowns during credit crunches.
- Estimation Risks: With nearly 80% right-censoring, tail estimation carries uncertainty. Sensitivity testing on ρ indicates that every 0.05 increment in correlation drastically widens tail losses.

1. Cell 2: Raw Data Profiling & Initial Inspection

- **Decision Taken:** Loaded the raw `loan_tape.csv` file to inspect overall dimensions, missing value (NaN) distributions, and data types.
- **Rationale:** To verify whether the empirical data adhered to the contractual schema and to detect data quality defects at origination.
- **Outcome:**
 - Identified 308,222 rows and 12 columns.
 - Detected missing values across critical fields: `origination_date` (6,078 missing), `maturity_date` (6,075 missing), `amortization_type` (1,208 missing), and `interest_rate` (3,049 missing).
 - High missingness in `default_date` (266,611 null) and `prepayment_date` (270,218 null) confirmed extensive right-censoring rather than corrupted entries.

2. Cell 5: Standardizing Date Formats

- **Decision Taken:** Converted date strings (`origination_date`, `maturity_date`, `default_date`, `prepayment_date`, `obs_end_date`) into standard `datetime` objects.
- **Rationale:** To enable duration calculations, date-difference operations, and temporal validation logic.
- **Outcome:** Standardized timestamp parsing and flagged unparseable temporal records for downstream cleaning.

3. Cell 6: Scanning Logical Inconsistencies & Categorical Defects

- **Decision Taken:** Audited for competing terminal events, events recorded prior to loan origination, and invalid categorical inputs.

- **Rationale:** A loan cannot default prior to being originated; furthermore, early prepayment and default represent mutually exclusive competing absorbing states.
- **Outcome:**
 - Identified 7,981 loans recording both a default and a prepayment date.
 - Identified 5,836 default dates and 4,286 prepayment dates occurring before loan origination.
 - Confirmed that 10,317 missing origination dates were mathematically recoverable via maturity date and contractual term.
 - Detected typos in `amortization_type` (ITALIAN, Frnch, bullet, XYZ) and implausible interest rates ($\$r \leq -0.02\%$ or $\$r = 1.50\%$).

4. Cell 7: Data Cleaning Pipeline & Data-Quality Log Construction

- **Decision Taken:** Executed deterministic cleaning rules and tracked all modifications inside an auditable `dq_log` dataframe:
 - Recovered 10,317 origination dates using:

$$\text{\texttt{\$origination_date}} = \text{\texttt{\$maturity_date}} - \text{\texttt{\$term (months)}}$$
 - Dropped 384 records with unrecoverable origination dates.
 - Set impossible pre-origination default (6,053) and prepayment (4,468) dates to `NaN`.
 - Standardized valid amortization structures (bullet, italian, french) and dropped unmapped XYZ/null entries (3,654 records).
 - Removed 7,572 rows with interest rates outside $$(0, 40\%]\$$.
- **Rationale:** To fulfill the Part 1 mandate of creating a reproducible, auditable Data-Quality Log without discarding salvageable records.
- **Outcome:** Retained 296,612 records (96.23% of the raw tape) and generated a structured audit table.

5. Cell 11: Amortization Schedule & Outstanding Balance (EAD) Engine

- **Decision Taken:** Implemented the `compute_ead()` routine to derive outstanding contractual principal at observation close (`obs_end_date`):
 - **Bullet:** Full principal maintained until maturity.
 - **Italian:** Linear principal repayment schedule.
 - **French:** Constant-annuity amortization formula.
- **Rationale:** To model Exposure at Default (EAD) based on contractual cash-flow depletion.
- **Outcome:** Generated valid `ead_obs_end` balances across obligors, revealing negative amounts and corporate-scale outliers exceeding €50M.

6. Cell 12: Outlier Truncation & Event State Labeling

- **Decision Taken:**

- Filtered out 8,793 records with invalid balances ($\text{issue_amount} \leq 0$ or $> \text{€}100,000$).
 - Assigned terminal absorbing states: **default**, **prepayment**, or **censored**.
- **Rationale:** To isolate the retail portfolio boundary and explicitly represent right-censoring alongside competing risks.
- **Outcome:**
 - Yielded a clean baseline sample of 287,819 retail loans.
 - Terminal state breakdown: 79.54% Censored (228,929), 10.88% Prepayment (31,312), and 9.58% Default (27,578).

7. Cell 13: Marginal Survival Analysis & Competing Risks (Part 2)

- **Decision Taken:** Formulated survival duration as loan tenure in months ($\tau = \text{end_date} - \text{origination_date}$) and fitted three parametric hazard distributions via right-censored Maximum Likelihood Estimation:

$$\ell(\theta) = \sum_i \left[e_i \log f(t_i; \theta) + (1 - e_i) \log S(t_i; \theta) \right]$$
- **Rationale:** To model seasoning hazards while accounting for right-censoring, selecting the optimal model via information criteria (AIC/BIC).
- **Outcome:**
 - **Log-Normal:** $s = 1.937$, $\text{scale} = 431.49$ | **AIC: 65,576.5** (Optimal fit selected)
 - **Weibull:** $c = 0.970$, $\text{scale} = 395.27$ | AIC: 66,030.5
 - **Exponential:** $\text{scale} = 371.48$ | AIC: 66,036.3
 - Confirmed that the non-monotonic Log-Normal hazard effectively captures consumer seasoning (initial hazard increase followed by gradual decay).

8. Cell 14: Latent Factor Copula & Asset Correlation (Part 3)

- **Decision Taken:** Formulated a single-factor Vasicek / Gaussian copula model, optimized ρ against empirical vintage default rates via bounded MLE, and computed the lower tail dependence coefficient (λ_L) for a Student- t ($\nu=4$) structure:

$$X_i = \sqrt{\rho} Z + \sqrt{1 - \rho} \epsilon_i$$
- **Rationale:** To capture correlated systemic defaults driven by latent macro shocks and test tail-dependence properties.
- **Outcome:**
 - Unconditional portfolio default rate: $\bar{p} = 0.0958$.
 - Estimated asset correlation: hit the optimization boundary at $\rho = 0.5000$.
 - Proved that while the Gaussian copula assumes asymptotic independence ($\lambda_L = 0$), a Student- t structure yields $\lambda_L = 0.2532$, indicating severe joint-loss clustering under economic downturns.

9. Cell 15: Regional Concentration & Monte Carlo Portfolio Loss Engine (Parts 4 & 5)

- **Decision Taken:**
 - Segmented portfolio exposure and baseline defaults by geographic region.
 - Simulated $N = 10,000$ portfolio loss paths over a 12-month horizon with 45% unsecured LGD across two specifications: **Independent Benchmark ($\rho = 0$)** versus **Fitted Copula ($\rho = 0.50$)**.
- **Rationale:** To evaluate the originator's geographic diversification claims and quantify the economic capital shortfall under assumed independence.
- **Outcome:**
 - **Diversification Fallacy:** Default rates diverged (Northeast at 16.6% vs. Southeast at 6.2%), but common exposure to the systemic factor ($\rho = 0.50$) induced concurrent failure across regions.
 - **Capital & Loss Metrics:**
 - **Expected Loss (EL):** €66,012 (Independent) vs. €66,670 (Copula).
 - **99.9% Value at Risk (VaR):** €112,907 (Independent) vs. **€1,411,392** (Copula).
 - **Economic Capital ($\text{VaR}_{99.9} - \text{EL}$):** €46,894 vs. **€1,344,722**.
 - **Final Recommendation:** Assuming independence understates required tail capital by **28.7x**; a minimum **32% Credit Enhancement** is required to safeguard senior tranches.

Rubric Criteria & Execution Alignment

- **Part 1 (10 Pts) — Data Cleaning & EDA:**
 - *Requirement:* Classify censoring without treating it as error, establish an auditable Data-Quality Log, and compute correct EAD schedules.
 - *Implementation:* Cell 7 **dq_log** table, invalid chronological date imputation, and Cell 11 closed-form amortization engines.
- **Part 2 (15 Pts) — Marginals & Competing Risks:**
 - *Requirement:* Separate prepayment from default hazards, benchmark $\geq 3\%$ parametric families via AIC, and evaluate seasoning.
 - *Implementation:* Cell 13 parametric MLE fit ranking Log-Normal (AIC: 65,576.5) above Weibull and Exponential.
- **Part 3 (40 Pts) — Copula & Latent Structure:**
 - *Requirement:* Justify latent factor architecture with economic reasoning and demonstrate tail-dependence dynamics.
 - *Implementation:* Cell 14 Vasicek latent factor engine, Gaussian $\lambda_L = 0$ proof, and Student- t ($\lambda_L = 0.2532$) tail comparison.
- **Part 4 & 5 (30 Pts) — Segmentation & Simulation:**
 - *Requirement:* Assess cross-segment correlation, compute EL/VaR/ES with MC simulation, and benchmark against independence.

- *Implementation:* Cell 15 10,000-path vectorized engine demonstrating the 28.7x tail capital shortfall under independence.
- **Part 6 (5 Pts + Executive Credibility) — Model Risk & Critique:**
 - *Requirement:* Candidly diagnose estimation boundaries and articulate where the model breaks.
 - *Implementation:* Explicitly documenting that $\rho = 0.50$ hit the upper boundary because vintage-level aggregation conflated portfolio seasoning and right-censoring with systemic macro shocks.

Recommended Closing Statement for Your Oral Exam

We do not accept the 99.9% VaR tail output at face value to recommend an uncritical capital reserve. Our post-estimation diagnosis identified that the asset correlation hit the 0.50 optimization bound because vintage-level grouping conflated portfolio seasoning and right-censoring with systemic macro shocks. Rather than treating this as a through-the-cycle baseline, we interpret the 28.7x capital divergence as an extreme macro-stress scenario that exposes the structural flaws of assuming independent defaults.